COM4115 / COM6115: Text Processing (2018/19)
Assignment: Text Compression

NB: *COM3110 students should* **NOT** *attempt this assignment – it is not part of the assessment for that module. It is for students of COM4115 / COM6115 only.*

Task in brief: To implement the Huffman Coding text compression algorithm in Python and to investigate the advantages and disadavantages of character-based vs word-based encoding.

Submission: Submit your assignment work *electronically* via MOLE. Precise instructions for what files to submit are given later in this document. Please check you have access to the relevant MOLE unit (i.e. `COM4115 Text Processing (AUTUMN 2018-19)` or `COM6115 Text Processing (AUTUMN 2018-19)`) and contact the module lecturer if not.

SUBMISSION DEADLINE: 12 noon, Monday week 11 (3 December, 2018)

Penalties: Standard departmental penalties apply for late hand-in and use of unfair means. *In particular you* **MUST NOT** *use anyone else's code for this assignment or provide your code to anyone else, regardless of whether it is open source or not. Huffman coding is a well-known technique and there are implementations of it on the Web. We are aware of them and your code will be checked against them. Your code will also be checked against that of your classmates. Last year around 30 students on this module failed this assignment due to copying code and/or received a disciplinary warning. Start in plenty of time and if you have problems with the assignment seek help from the lab demonstrators or your module lecturer.*

Materials Provided

Data file: The Project Gutenberg version of Melville's novel `mobydick.txt` is supplied as a data file for development and testing (you may want to use a small portion of this text for initial development and testing).

Test Harness: A python script called `test-harness.py` is supplied which simulates the tests that will be run on your code when submitted. It won't actually run all the tests we will run, but it will attempt to execute the programs you are asked to submit using the program names and input file, output file and command line flag conventions you are asked to follow in the assignment. This script will print error messages if it detects any obvious problems in finding your programs or in executing them and will print a success message if everything appears to be present and named appropriately. *You must ensure that your code can be successfully run by this test harness prior to submission . If your code fails the* `test-harness.py` *test then it will not be run and you will lose all the marks allocated to implementation – see Assessment Criteria below.*

Task Description

Your task is write a two Python programs, one called `huff-compress.py` and the other called `huff-decompress.py`. For simplicity both programs and all their input and output files should be assumed to be in the same directory, which is also assumed to be the current working directory.

`huff-compress.py` should take a UTF-8 encoded plain text file (we will refer to this as `infile.txt`) as its first argument and should behave as follows:

1. Using `infile.txt` as data, the program should count the frequency of the symbols to be assigned Huffman codes (either characters or words – see Note 1 below) and should build a zero order model for symbol probabilities. To estimate the probabilities in the zero order model use the relative frequency of each symbol in `infile.txt` (the relative frequency of a symbol s in a text T is the proportion of symbol occurrences in T that are s). Take special care in dealing with non-printable characters.

2. The program should compute a Huffman code-tree for your `infile.txt` symbol set, using the zero order model computed in the previous step to assign probabilities to the nodes in your tree. One method to do this is described below in Note 2.

3. Once the code tree is complete the program should write two output files (see note 3):

   (a) The first, which should be called `infile-symbol-model.pkl`, should contain a serialized version of your symbol model (which is probably a dictionary), produced using the Python Pickle module. See Note 4 below.

   (b) The second, which should be called `infile.bin`, should contain the text of `infile.txt` compressed using your symbol model. Be careful to write to write this data out as binary data and not, e.g., as an ASCII representation of 1s and 0s. See Note 5 below. Also, to deal with the problem of your encoding potentially ending in the middle of a byte, be sure to add a pseudo-EOF (end-of-file) character to your symbol model and to add this character to the end of your encoded data. See Note 6.

`huff-decompress.py` should take a binary file encoded using your huff-compress.py progam (we will refer to this as `infile.bin` as its first argument and should behave as follows:

1. It should look for a file called `infile-symbol-model.pkl` in the same directory as `infile.bin` and should use `pickle` (Note 4) to read in and recreate the symbol model.

2. Using the symbol model acquired in the previous step the program should read the binary data file `infile.bin`, decode it into plain text and then write it to an output file called `infile-decompressed.txt`. Obviously, `infile-decompressed.txt` should match the original `infile.txt` (the Unix/Linux utility `diff` can be used to test this).

   (Normally compression programs would simply write the output to a file called `infile.txt`, overwriting any existing file with that name. This is not a good idea during development or assessment of such a program, however.)

What to Submit

Your assignment work is to be submitted *electronically* using MOLE, and should include:

1. Your two Python programs `huff-compress.py` and `huff-decompress.py`. Do *NOT* submit any other Python codes files.

2. A short report (as a **pdf** file), which should *NOT EXCEED 2 PAGES IN LENGTH*. The report should include a brief description of the extent of the implementation achieved (this is especially important if you have not completed a sufficient implementation for performance testing). It should also address the following questions:

    1. How does the compression performance of character-based versus word-based Huffman compare? Specifically:

        (a) How big (in bytes) are the compressed text files produced in each case?
        (b) How big (in bytes) are the symbol models produced in each case?
        (c) How long does it take to
            i. build the symbol model?
            ii. encode the input file given the symbol model?
            iii. decode the compressed file? (including unpickling the symbol model)
            See Note 7.

    2. How could various aspects of the performance, as identified in response to the previous question, be improved?

    Any other observations you believe are pertinent to comparing character-based and word-based Huffman coding should also be made. Graphs/tables may be used in presenting your results, to aid exposition.

Assessment Criteria

A total of 25 marks are available for the assignment and will be assigned based on the following criteria.

Implementation and Code Style (15 marks)
How fully and correctly has the functionality specified in the "Task Description" section above been implemented? How efficient is the implementation (i.e. how quickly are results returned)? Have appropriate Python constructs been used? Is the code comprehensible and clearly commented? You can expect your program will be run on `mobydick.txt` and perhaps on other texts as well. Amongst other things it will be checked that the result of compressing and then uncompressing a file is identical to the original file prior to compression.

Report (10 marks)
Is the report a clear and accurate description of the implementation? How complete and accurate are the answers to the questions posed above in the "What to Submit" section? What conclusions can be drawn about character-based versus word-based Huffman coding from these results?

Notes and Comments

1. **Tokenization**: Your program should allow a user to decide whether the symbols to be assigned Huffman codes will be characters or words. This should be specifiable via command line arguments which follow the syntax: `-s char` for generating codes at the symbol level and `-s word` for generating codes at the word level. To identify the words in the input text you should:

   1. treat every continuous sequence of alphabetic characters (of either upper or lower case) as a word – i.e. use a regular expression that matches "a-zA-Z" and consider each match a word;
   2. treat every other character (numerals, punctuation, white space) as a word.

2. **Building a Huffman Code Tree**:

   a. Create a class whose instances will serve as nodes in (binary-branching) code trees. Instances might have attributes for the symbol (for use in leaf nodes), the node's probability, and for storing left and right subtrees (all of which might be initialised to `None`).

   b. Create a list of (leaf) nodes for the individual symbols (each also storing the node's associated probability).

   c. Sort this list on the basis of node probability, so that the two lowest probability nodes can be found together at one end of the list.

   d. Remove the two lowest probability nodes from the list, and combine them to form a new subtree, i.e. as the left and right daughters of a new node.

   e. Add the new node back into the list of nodes, maintaining sorted order (w.r.t. node probability). An easy way to do this is just to append it and then re-apply the sort method (which is not optimally efficient – you may want to consider a more efficient alternative).

   f. Repeat steps (d) and (e) until only a single node remains in the list.

   g. Finally, traverse the tree of the remaining node, to assign binary codes to the symbols and return them (e.g. as a dictionary).

3. **Format of Compressed File(s)**: A more user-friendly approach to implementing text compression using Huffman coding is to put both the symbol model and the compressed text into a single output file. However this requires more effort than is warranted for this assignment and makes it harder to independently assess the size of the symbol model and compressed text files for different choices of symbol, i.e. for word-based or character-based encodings.

4. **Pickling your Data**: `pickle` is a Python module for serializing data, allowing complex data structures, such as dictionaries, to be written to and read from disk (i.e. they may be "preserved"). In the current case after `importing pickle` you should be able get away with using just the `pickle.dump` and `pickle.load` methods. See `https://docs.python.org/3/library/pickle.html?highlight=pickle#module-pickle` for full details. You might find `https://wiki.python.org/moin/UsingPickle` and `https://pythontips.com/2013/08/02/what-is-pickle-in-python/` helpful as well.

5. **Writing and Reading Binary Data**: Consider the example of Huffman coding from the first lecture on Text Compression where the string 'eefggfed' is coded as in binary as: 10101101111111101001. If you naively write this to disk as a string each binary symbol ('0' or '1') will be written in one byte resulting in a string of length 20 being used to represent an original string of length 8! This is hardly text compression. To get around this problem your string representation of a binary number must be converted to a binary representation where each 8 digits in the string representation (which at one byte each takes up $8 \times 8 = 64$ bits) are converted to into a single byte (8 bits). To do this Python you should use the `int` built-in which if you call it as

```
c = '11001100'
b = int(c,2)
```

will take an 8-character string such as `c` and convert into an integer base 2 representation (i.e. one byte).

Thus you should think of buidling up the whole text to be encoded into a single string of '1's and '0's and then loop through this 8 characters at a time converting each 8 characters into a one byte binary representation. If these one byte representations are added to an `array` then a proper binary representation of your encoding may be constructed. Here are some code fragments to illiustrate this:

```
import array
...
codearray = array.array('B')  # create array of bytes
...
    # somewhere inside a for loop where 8 string 1's and 0's are being
    # placed into c
    codearray.append(int(c, 2)) # convert c to binary and add to array
    ...

# Write the array at once to a file
#
f = open('workfile', 'wb')
codearray.tofile(f)
f.close()
```

You may find this discussion helpful: `https://www.linuxquestions.org/questions/programming-9/writing-binary-data-under-python-718165/`. Also helpful, if you are developing on a Unix-based platform (Linux, Mac OSX), is the system command `hexdump`, which prints out a dump of a file's contents in hexadecimal, 8 bytes per line. The `-C` flag also prints the ASCII interpretation of the line. This lets you check, on small example files during development, that you are getting things right.

6. **Pseudo-EOF**: Your binary encoding of the input text may be an arbitrary number of bits in length. However, when you go to write it to file you will need to write an integer number of bytes. Returning to the example from the lecture mentioned above, the 8 character

original string 'eefggfed' will end up being encoded in 20 bits as 10101101111111101001. When this is converted to 8-bit bytes we end up with two bytes + 4 bits. To write the last 4 bits we need to add 4 more bits to make up an integer number of bytes. The problem is if we just add arbtirary bits to finish off the last byte then when we come to decode our decoder may interpret these as further characters to be decoded. How is it to know the difference?

There are a number of solutions to this problem. The solution recommended here is to add a special symbol to your symbol set, call it ■, that serves as a pseudo-EOF marker. It gets used just once – it is the last symbol you add to your encoded output and anything can be placed in any residual bits in the last byte following it (e.g. '0's). On decoding, when ■ is encountered you know you are at the end of the text to be decoded. Further discussion of this approach can be found at `https://stackoverflow.com/questions/14667681/last-byte-in-huffman-compression` and in the pdf file at `https://web.stanford.edu/class/archive/cs/cs106b/cs106b.1126/handouts/220%20Huffman%20Encoding.pdf`.

7. **Timing**: There are various ways to time your code's execution. If you are on a Unix-based platform (Linux, Mac OSX) you can use the Unix system command `time`. There are Windows equivalents, but they are not so straightforward. Alternatively there are cross-platform timing commands available with Python, which you can use.